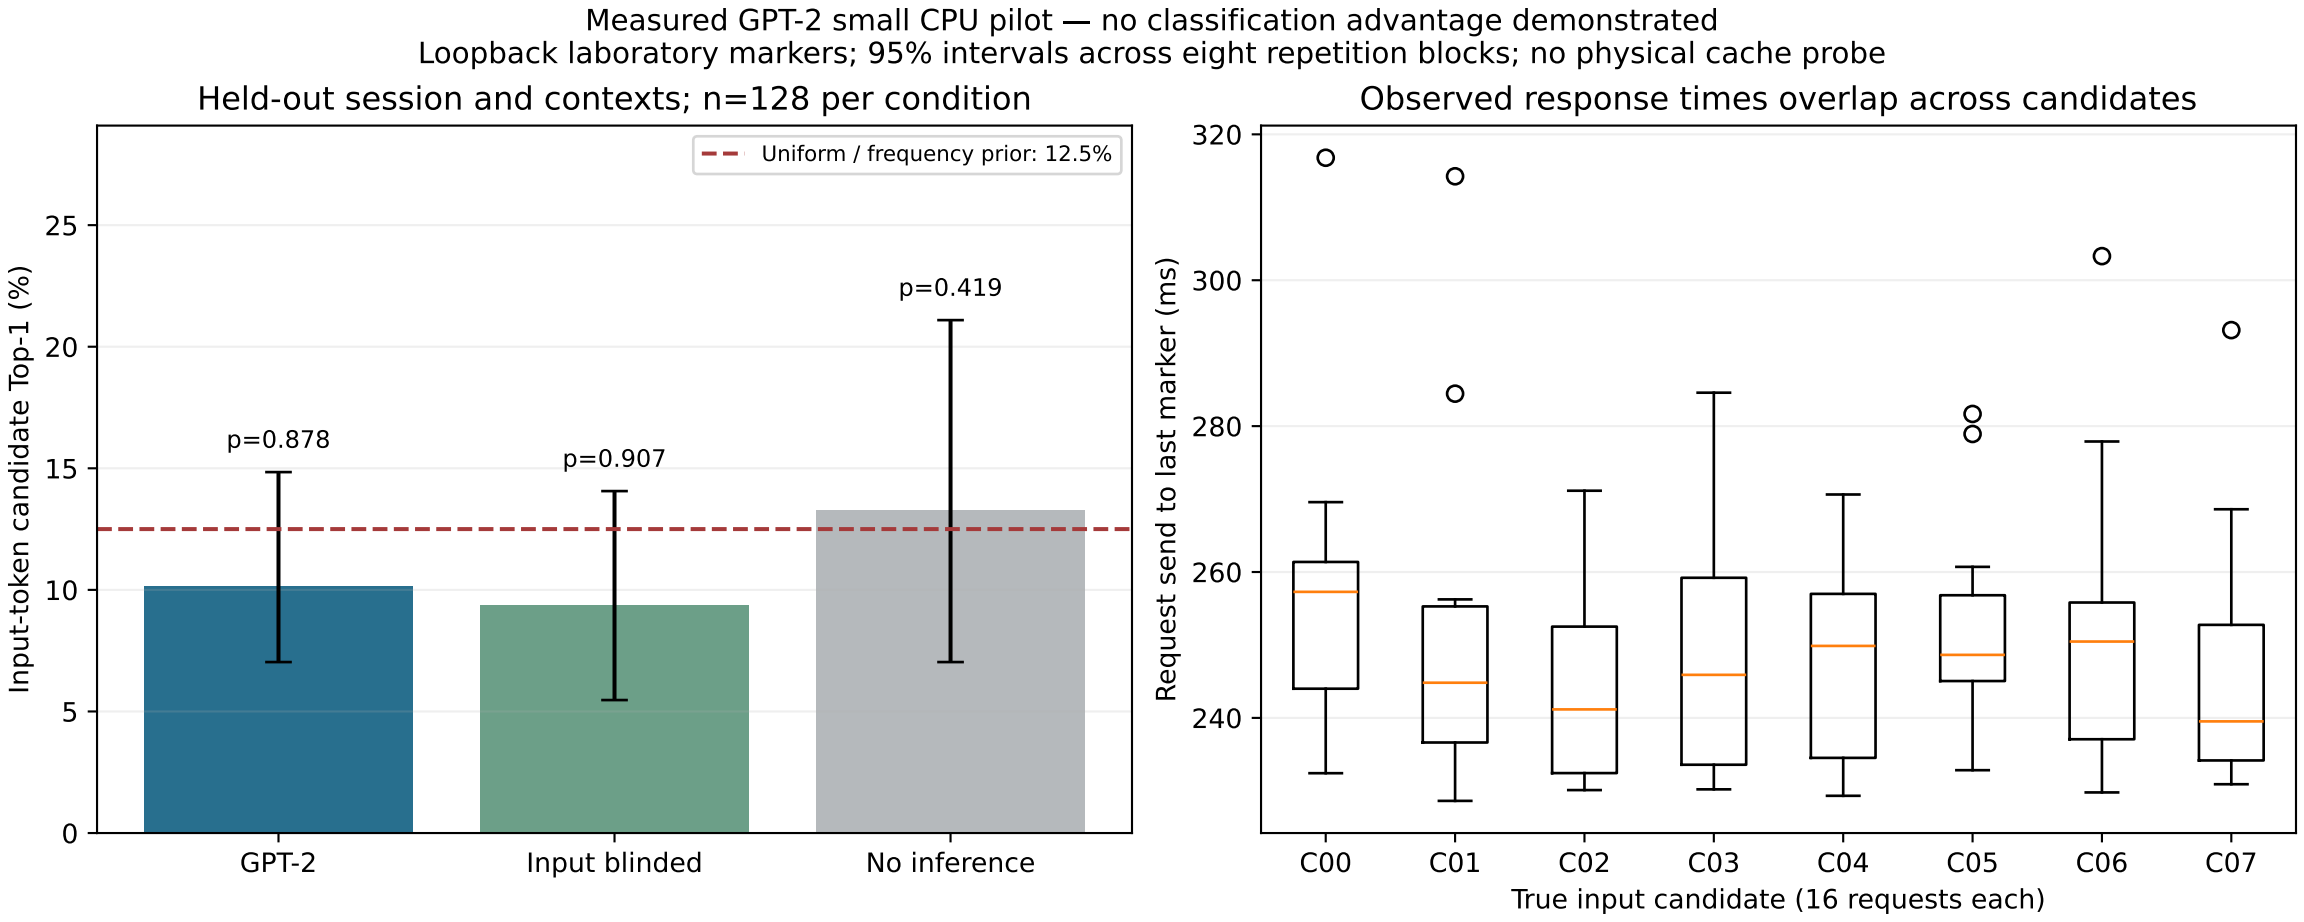
Measured GPT-2 small CPU pilot — no classification advantage demonstrated
Loopback laboratory markers; 95% intervals across eight repetition blocks; no physical cache probe
Held-out session and contexts; n=128 per condition
Uniform / frequency prior: 12.5%
Input-token candidate Top-1 (%)
0
5
10
15
20
25
p=0.878
p=0.907
p=0.419
GPT-2
Input blinded
No inference
Observed response times overlap across candidates
Request send to last marker (ms)
240
260
280
300
320
C00
C01
C02
C03
C04
C05
C06
C07
True input candidate (16 requests each)